

Descriptive statistics/plots continued  
and data transformations

# Overview

## Review

- Plots and statistics of categorical data
- Measures of central tendency

## Continuation of statistics and plots of quantitative data

- Measures of spread
- Two quantitative variables

Start on data transformations using dplyr

# Reminder: Homework 2

It is due on Gradescope by 11pm on Monday July 14<sup>th</sup>

- Question 4 involves reading a short article and commented on it, so you can get started on this right away

How is the homework going so far?

# Review: Categorical data

Categorical variables take on one of a fixed number of possible values

For categorical variables we usually want to view:

- **Frequency table:** How many items are each category or
- **Relative frequency table:** The proportion (or percentage) of items in each category

# Vector of drinking behavior

```
> drinking_vec <- profiles$drinks
```

# Frequency and relative frequency tables

```
> drinks_table <- table(drinking_vec)
```

```
> prop.table(drinks_table)
```

	age	body_type	diet	drinks	drugs	education
1	22	a little extra	strictly anything	socially	never	working on college/university
2	35	average	mostly other	often	sometimes	working on space camp
3	38	thin	anything	socially	NA	graduated from masters program
4	23	thin	vegetarian	socially	NA	working on college/university
5	29	athletic	NA	socially	never	graduated from college/university
6	29	average	mostly anything	socially	NA	graduated from college/university

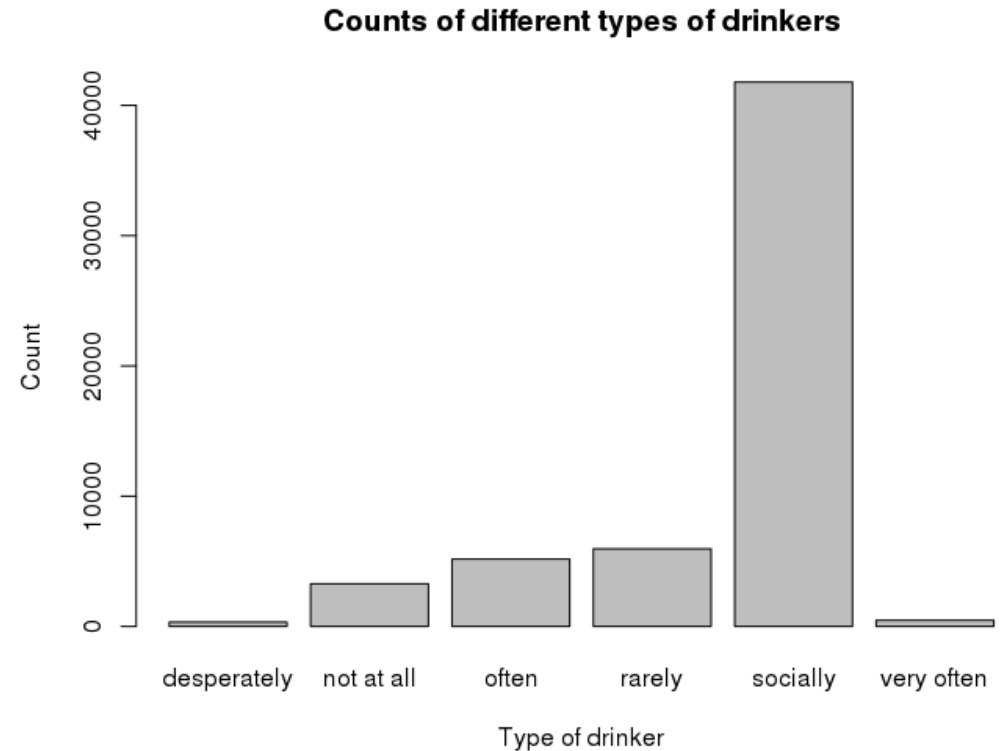
# Review: Visualizing categorical data

We can plot the number of items in each category using a bar plot

```
barplot(drinks_table,  
        ylab = "Count",  
        xlab = "Type of drinker")
```

We can also use the `pie()` function to create pie charts

```
pie(drinks_table)
```

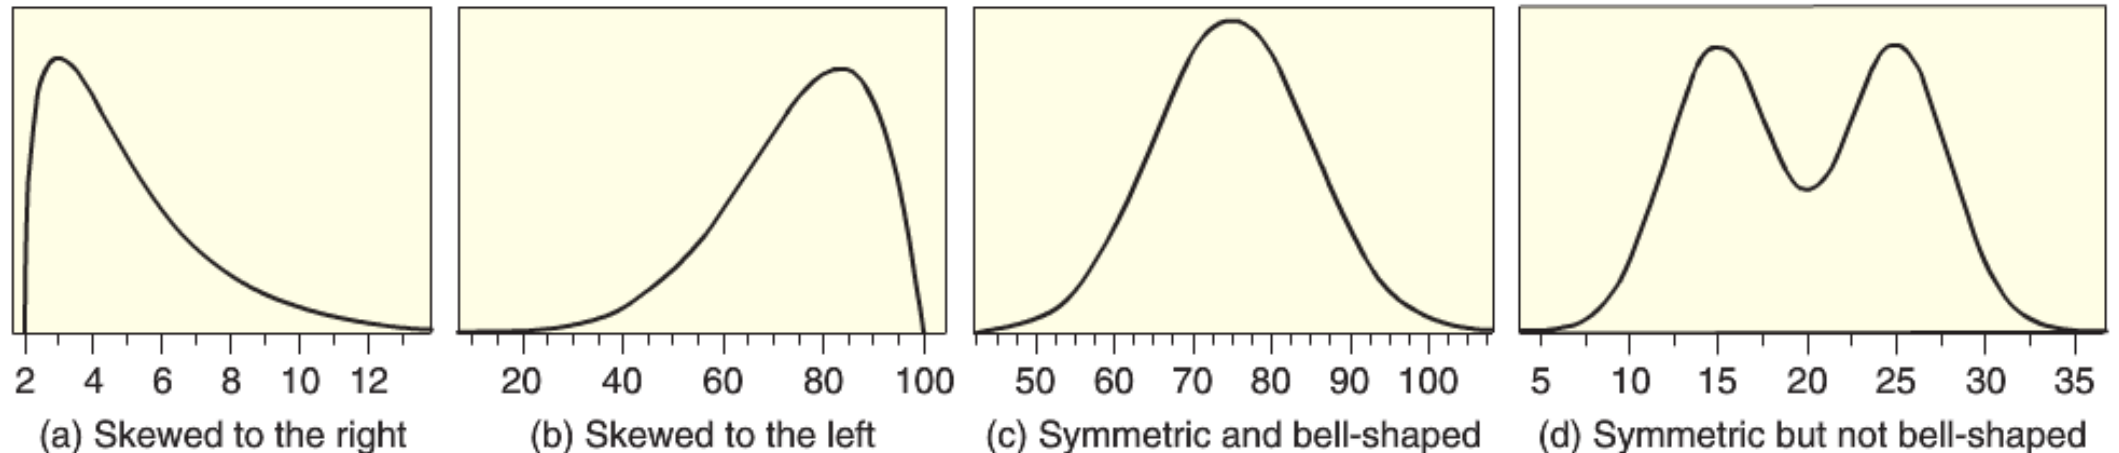


# Review Visualizing quantitative data

We can visualize quantitative data using histograms

```
hist(profiles$height, breaks = 50)
```

Common shapes of histograms are:



# Review Measures of central tendency

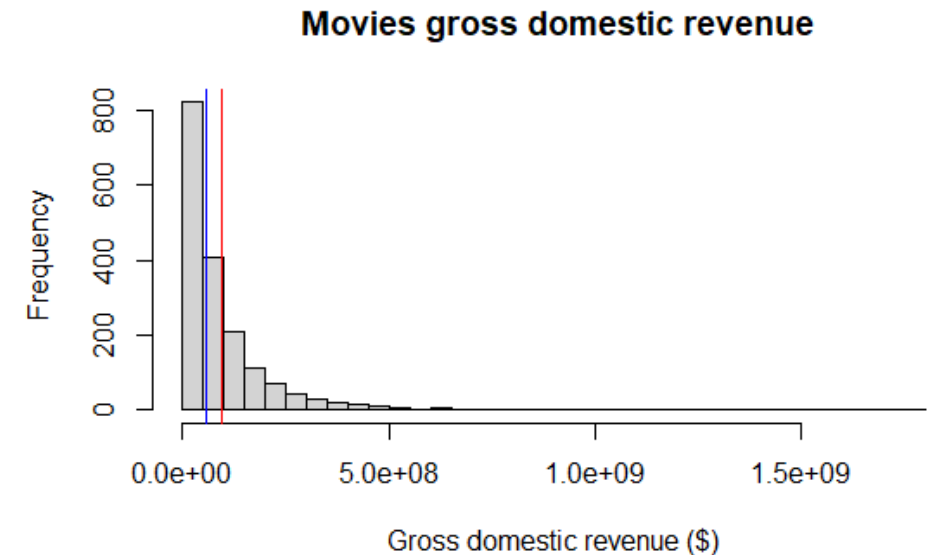
One common measure of central tendency is the **mean**

$$\text{mean}(x, \text{na.rm} = \text{TRUE}) \quad \frac{1}{n} \sum_{i=1}^n x_i$$

The **median** is the value such that half of the data is less than the median and half are greater than the median

$$\text{median}(v, \text{na.rm} = \text{TRUE})$$

The median is resistant to extreme values while the mean is not



Example:

Mean US salary = \$72,641

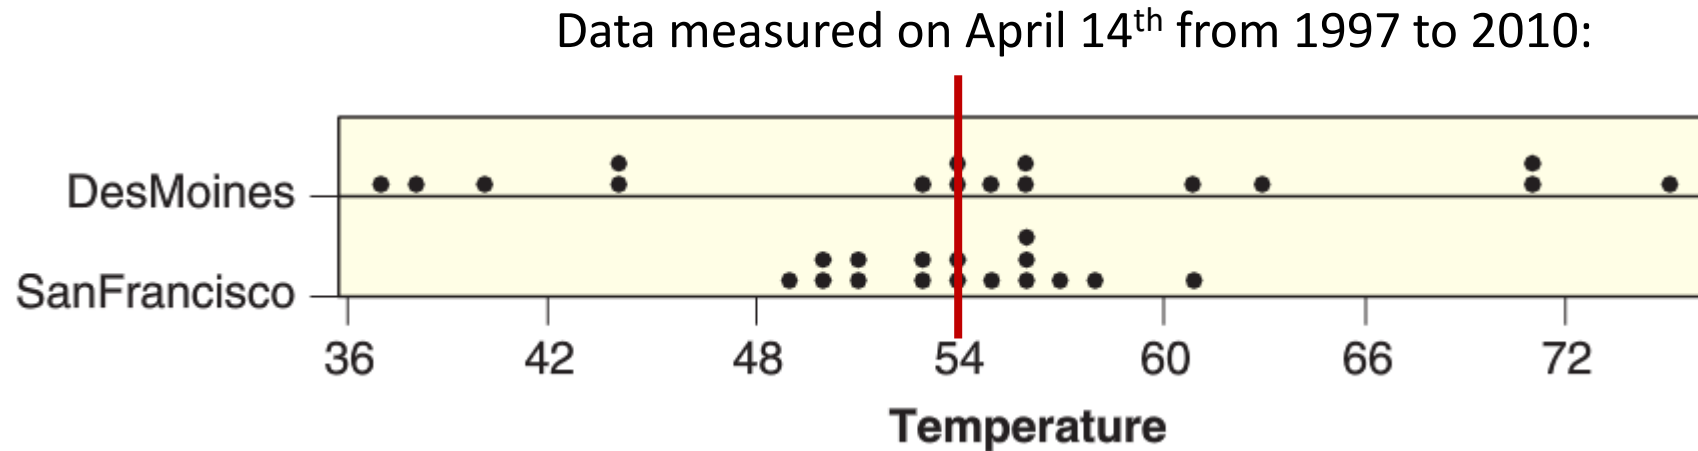
Median US salary = \$51,939

# Measures of spread



# Measure of spread 1: standard deviation

The **standard deviation** is a statistic that quantifies how far the data is spread



Mean temperature (°F):	Des Moines = 54.49	San Francisco = 54.01
Standard deviation (°F):	Des Moines = 11.73	San Francisco = 3.38

# Example: computing the standard deviation

Suppose we had a sample with  $n = 4$  points:

$$x_1 = 8, \quad x_2 = 2, \quad x_3 = 6, \quad x_4 = 4,$$

We can compute the mean using the formula:

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i = \frac{1}{4} \cdot (x_1 + x_2 + x_3 + x_4) = \frac{1}{4} \cdot (8 + 2 + 6 + 4) = 5$$

The standard deviation can be computed using the formula:

$$s = \sqrt{\frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2}$$

$$s = \sqrt{\frac{1}{4-1} \sum_{i=1}^n (x_i - 5)^2}$$

(remember order of operations!)

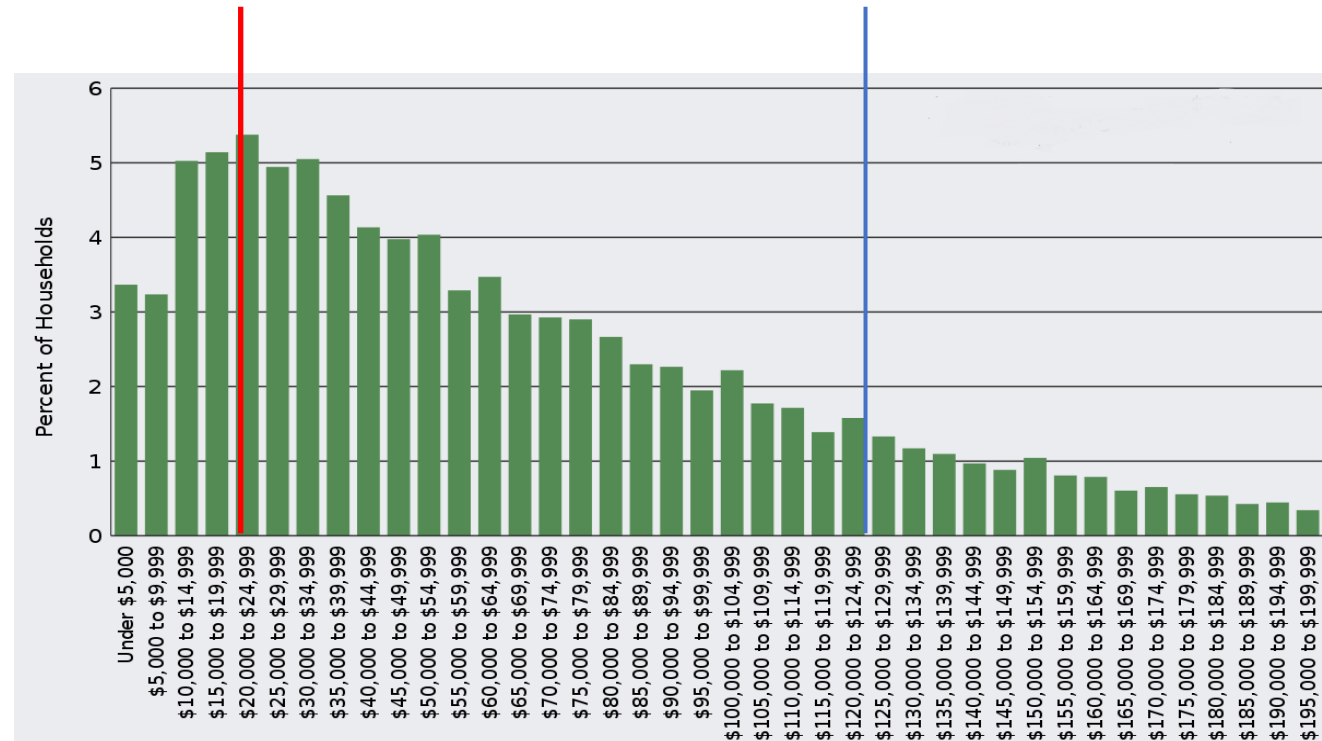
# Percentiles

The **P<sup>th</sup> percentile** is the value of a quantitative variable which is greater than P percent of the data

For the US income distribution what are the 20<sup>th</sup> and 80<sup>th</sup> percentiles?

20<sup>th</sup> percentile = \$21,430

80<sup>th</sup> percentile = \$112,254



R: `quantile(v, .95)`

# Five Number Summary

A **five-number summary** is a set of five descriptive statistics that provides a concise overview of a dataset's distribution

**Five Number Summary** = (minimum,  $Q_1$ , median,  $Q_3$ , maximum)

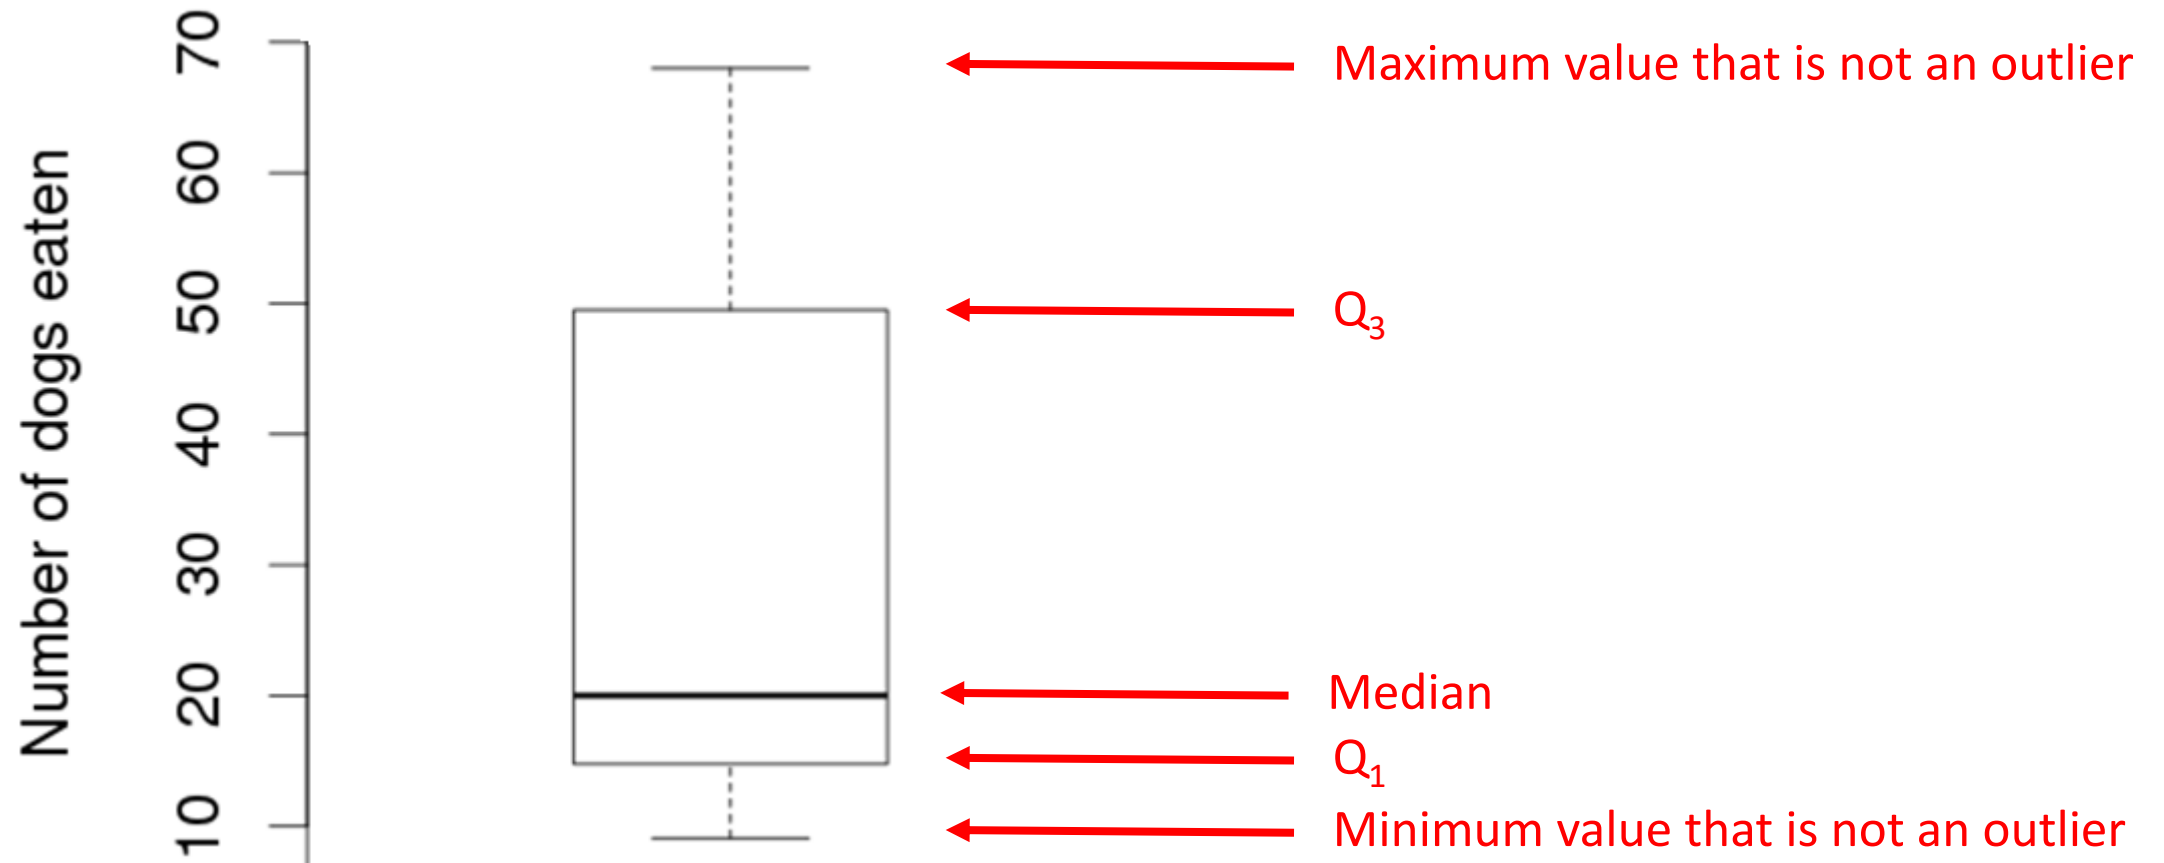
$Q_1$  = 25<sup>th</sup> percentile (also called 1<sup>st</sup> quartile)

$Q_3$  = 75<sup>th</sup> percentile (also called 3<sup>rd</sup> quartile)

Roughly divides the data into fourths

**Measure of spread 2: Interquartile range (IQR) =  $Q_3 - Q_1$**

# Box plots can also visualize quantitative data



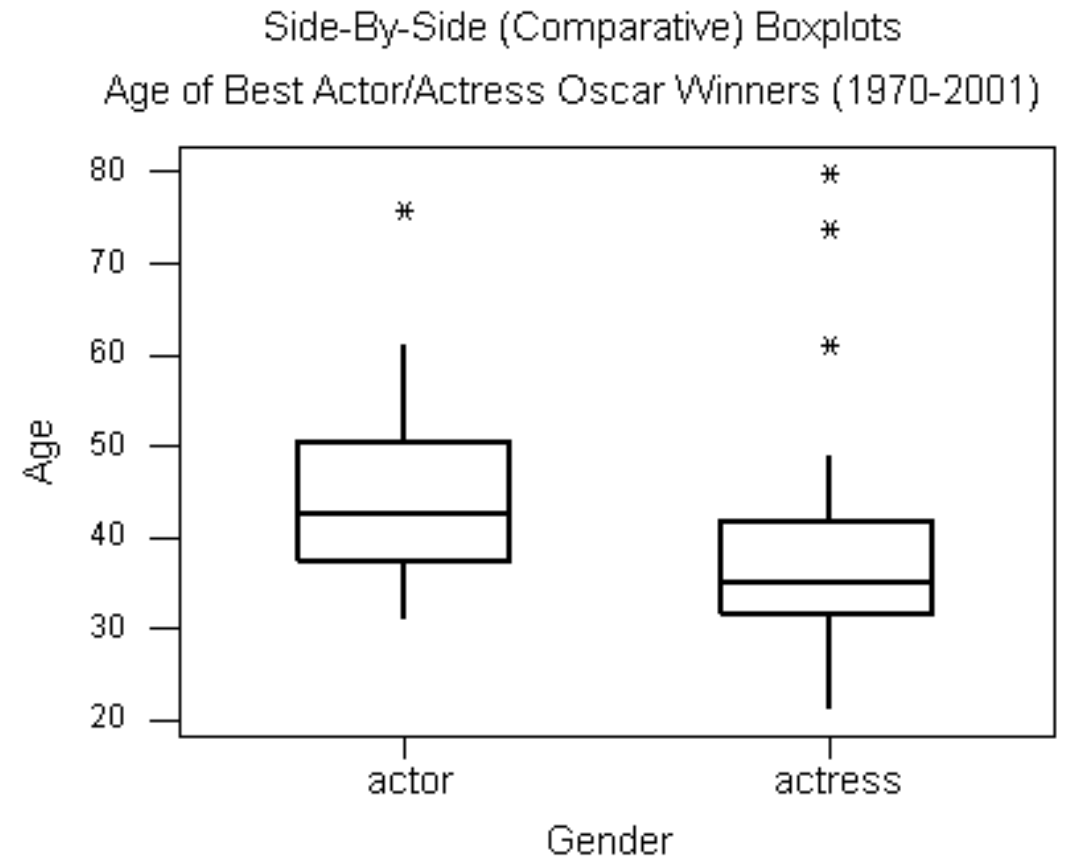
R: `boxplot(v)`

# Side-by-side boxplots

Boxplots are particularly useful for comparing distributions!

Let's look at the ages that people won the best actor/actress Oscar

What does this figure tell us?



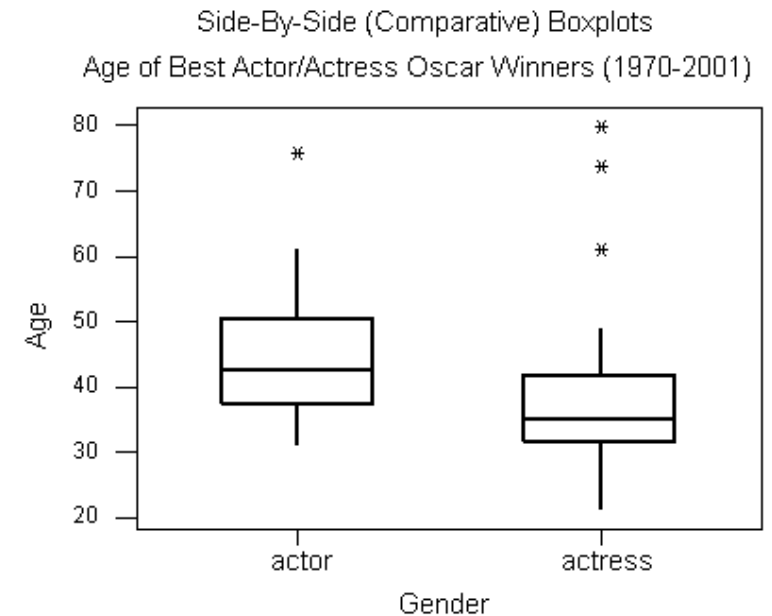
# Outliers

Outliers on boxplots are values that are more than  $1.5 * IQR$

What should we do if we have outliers?

Investigate:

- If there are due to an error, remove them
- If not, need to account for them



# Questions?



Let's try it in RStudio!



Visualizing two quantitative variables

# CitiBike data

Let's look at the bike share data from NYC

```
> load('daily_bike_totals.rda')
```



## CitiBike analysis

	date	trips	precipitation	snow_depth	snowfall	max_temperature	min_temperature
1	2013-07-01	16650	0.8385830	0	0	77.00	71.96
2	2013-07-02	22745	0.0787402	0	0	82.04	71.96
3	2013-07-03	21864	0.5314960	0	0	82.94	73.04
4	2013-07-04	22326	0.0000000	0	0	87.08	75.02
5	2013-07-05	21842	0.0000000	0	0	89.96	75.92
6	2013-07-06	20467	0.0000000	0	0	91.94	78.08
7	2013-07-07	20477	0.0000000	0	0	91.94	78.08
8	2013-07-08	21615	0.2204720	0	0	89.06	73.04

What does each case correspond to?

# Line plots

We can use the `plot(x, y)` function to create line plots

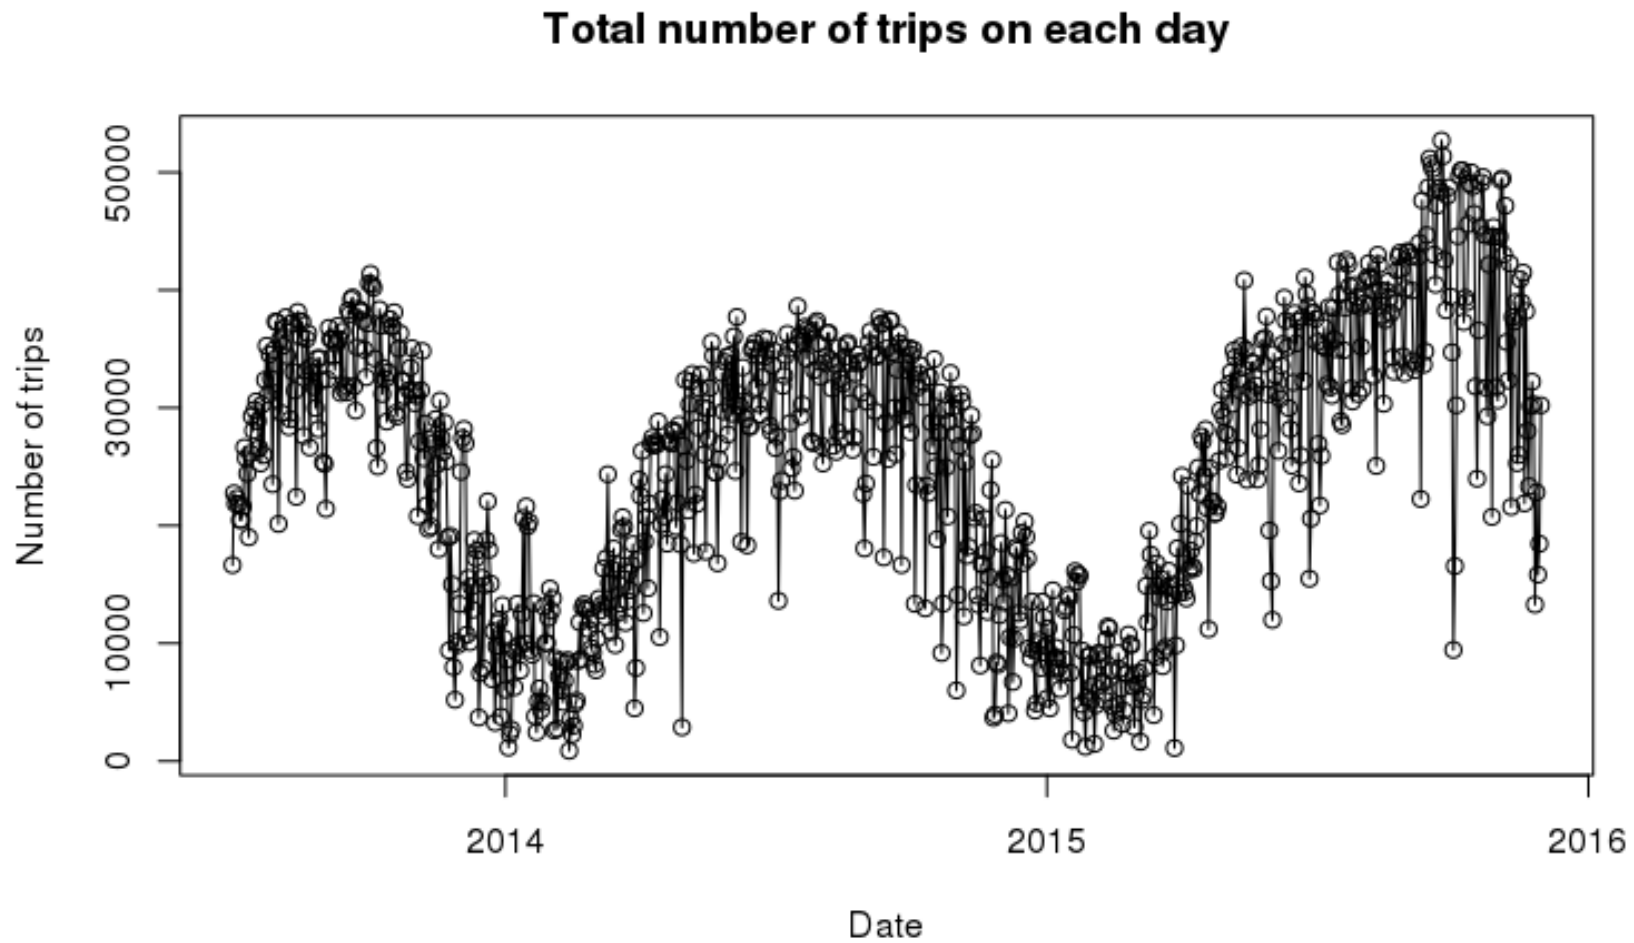
# we can connect the points in a plot using

```
plot(x, y, type = 'l') # line graph
```

```
plot(x, y, type = 'o') # both points and a line
```

```
plot(bike_daily_data$date, bike_daily_data$trips,  
     type = 'o',  
     xlab = "Date",  
     ylab = "Number of trips",  
     main = "Total number of trips on each day")
```

# Line plots

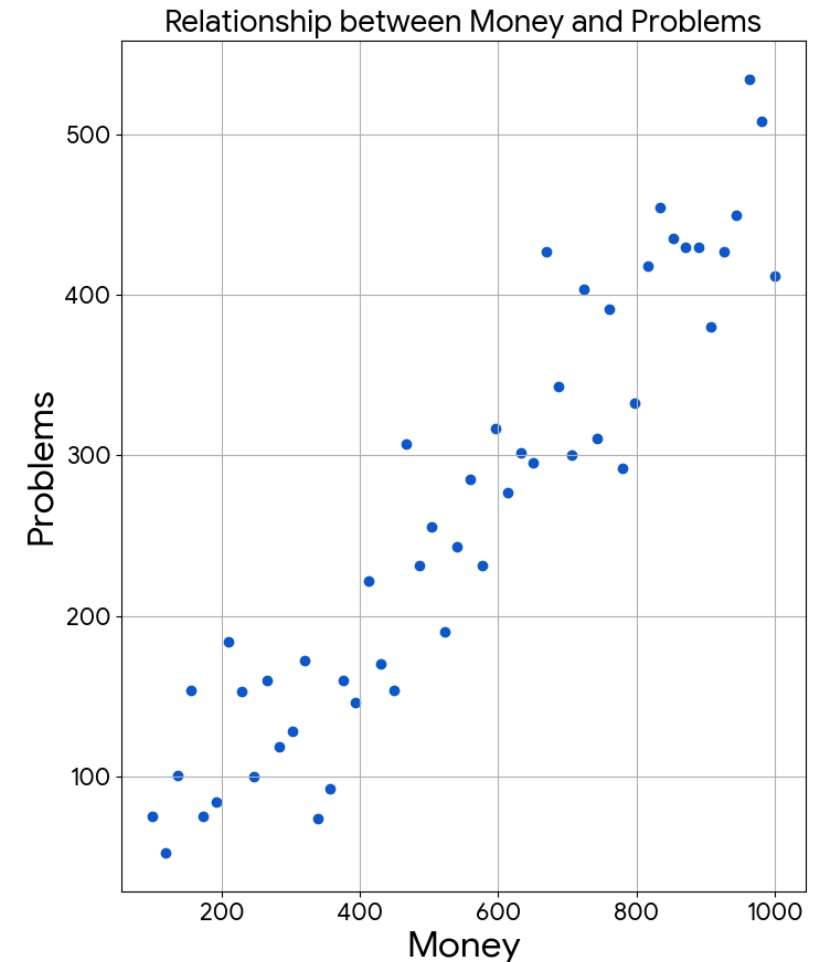


# Scatterplot

A **scatterplot** graphs the relationship between two variables

- Each axis represents the value of one variables
- Each point the plot shows the value for the two variables for a single data case

If there is an explanatory and response variable, then the explanatory variable is put on the x-axis and the response variable is put on the y-axis



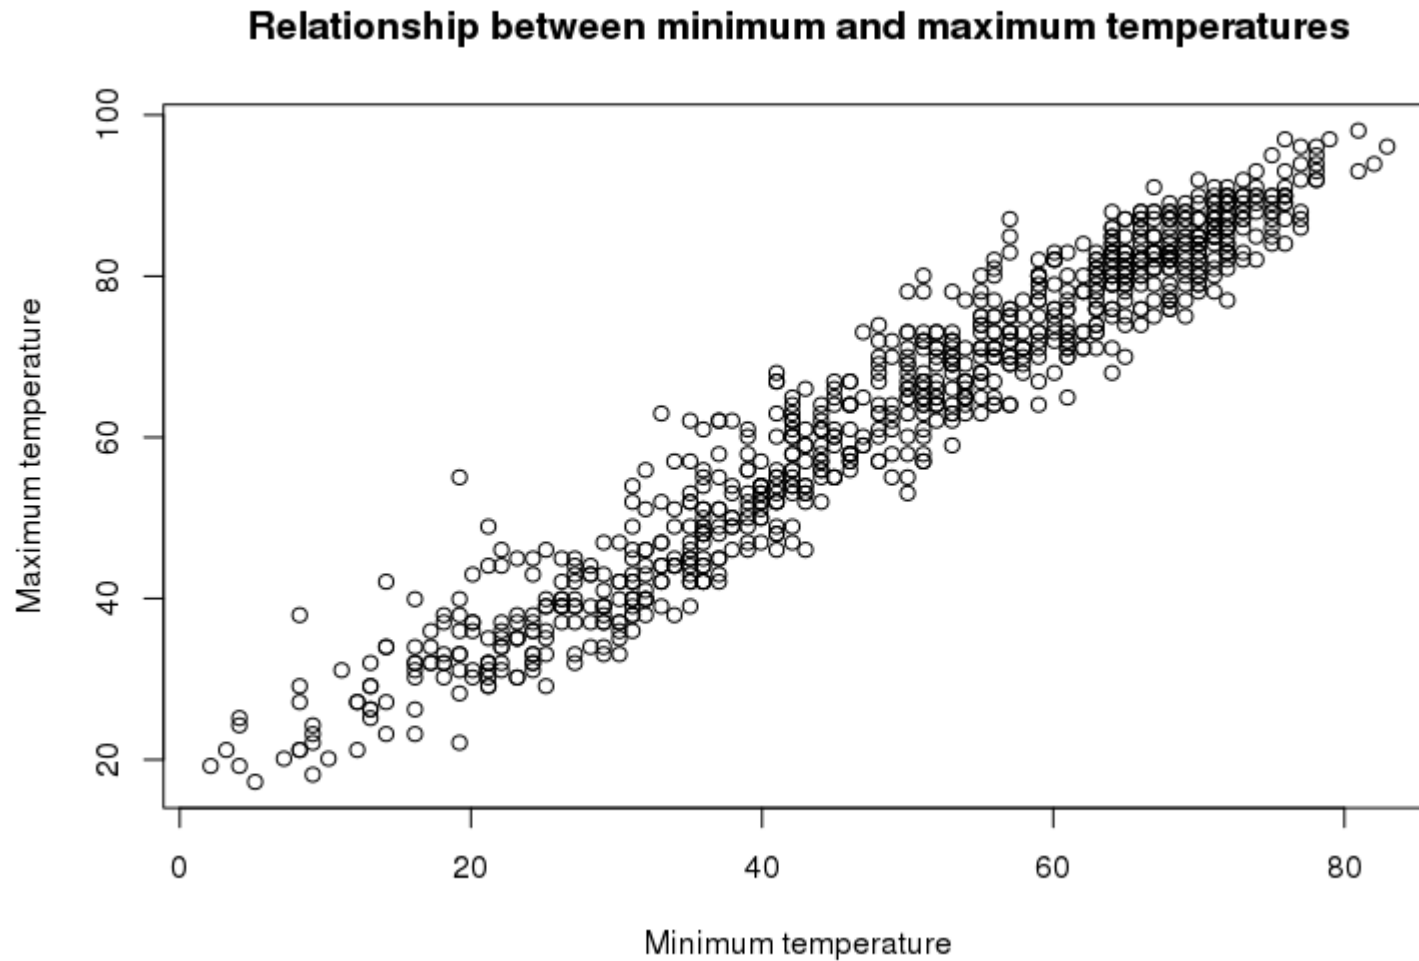
# Scatter plots

We can use the `plot(x, y)` function to create scatter plots

Can you create a scatter plot of the relationship between the minimum and maximum temperatures?

```
plot(bike_daily_data$min_temperature,  
     bike_daily_data$max_temperature,  
     xlab = "Minimum temperature",  
     ylab = "Maximum temperature",  
     main = "Relationship between min and temp")
```

# Scatter plots



# The correlation coefficient

The **correlation** is measure of the strength and direction of a linear association between two variables

$$r = \frac{1}{(n - 1)} \sum_{i=1}^n \left( \frac{x_i - \bar{x}}{s_x} \right) \left( \frac{y_i - \bar{y}}{s_y} \right)$$

R: `cor(x, y)`



# Properties of the correlation

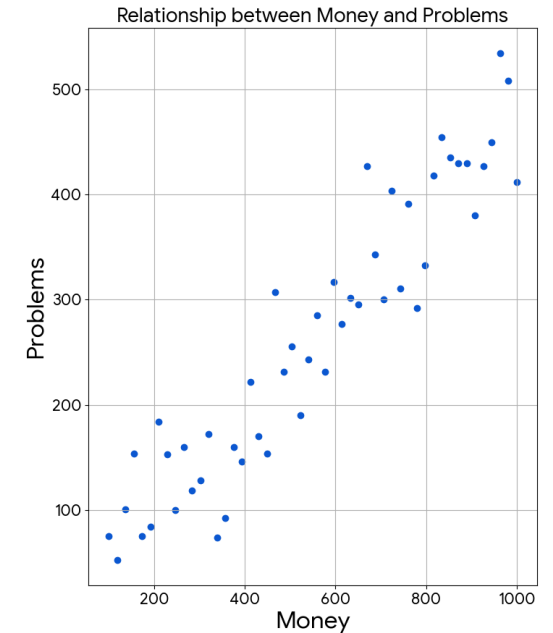
Correlation is always between -1 and 1:  $-1 \leq r \leq 1$

The sign of  $r$  indicates the direction of the association

Values close to  $\pm 1$  show strong linear relationships, values close to 0 show no linear relationship

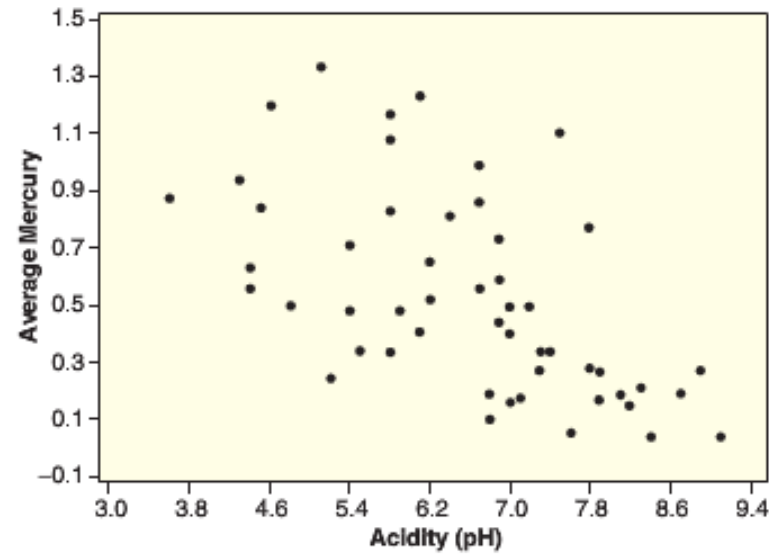
Correlation is symmetric:  $r = \text{cor}(x, y) = \text{cor}(y, x)$

$$r = \frac{1}{(n-1)} \sum_{i=1}^n \left( \frac{x_i - \bar{x}}{s_x} \right) \left( \frac{y_i - \bar{y}}{s_y} \right)$$

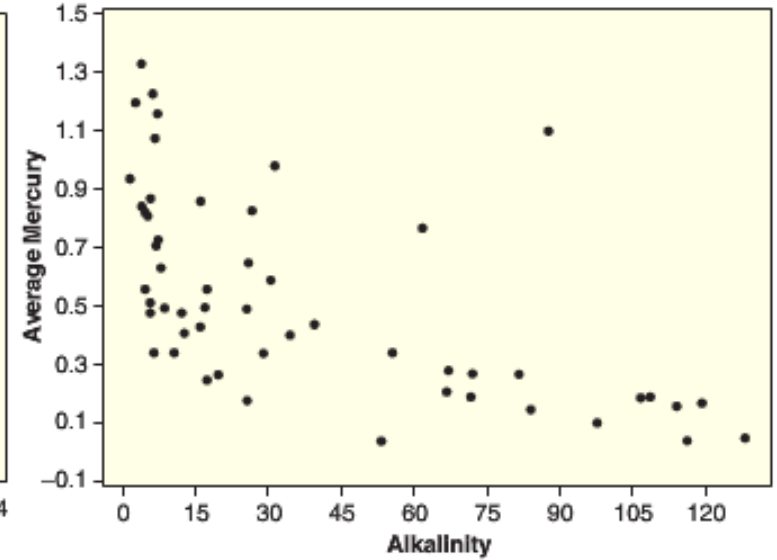


# Florida lakes

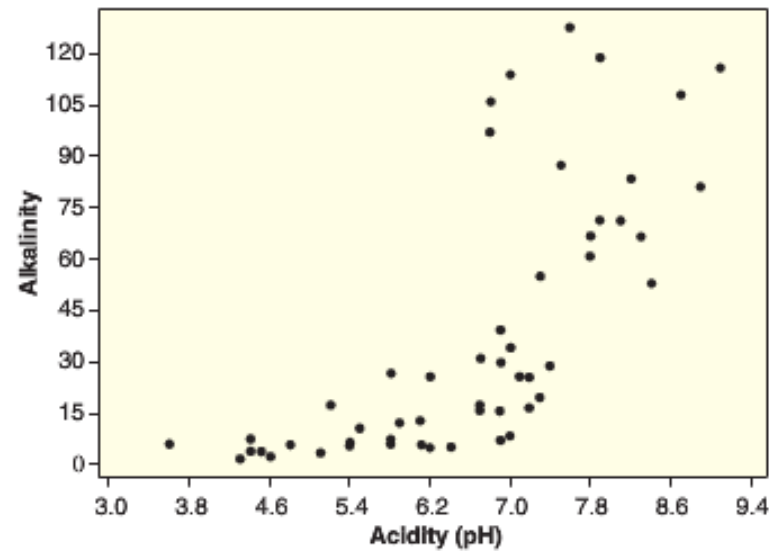
## Correlation game



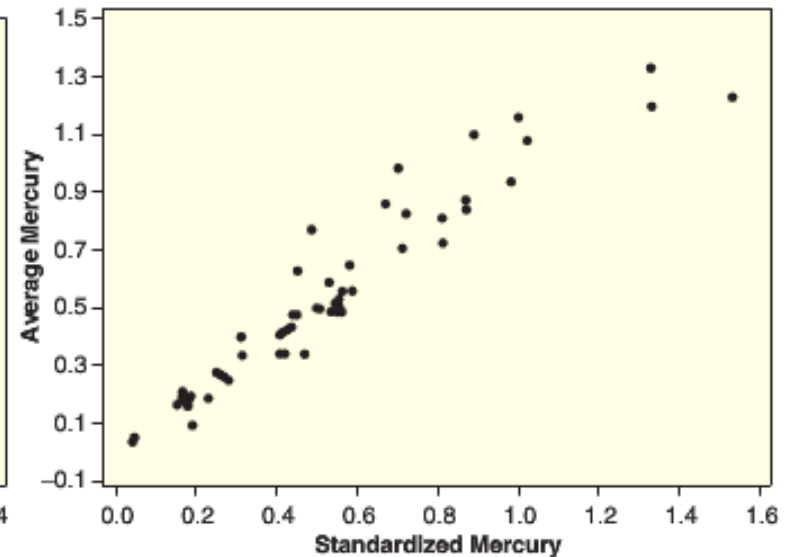
(a) Average mercury level vs acidity



(b) Average mercury level vs alkalinity



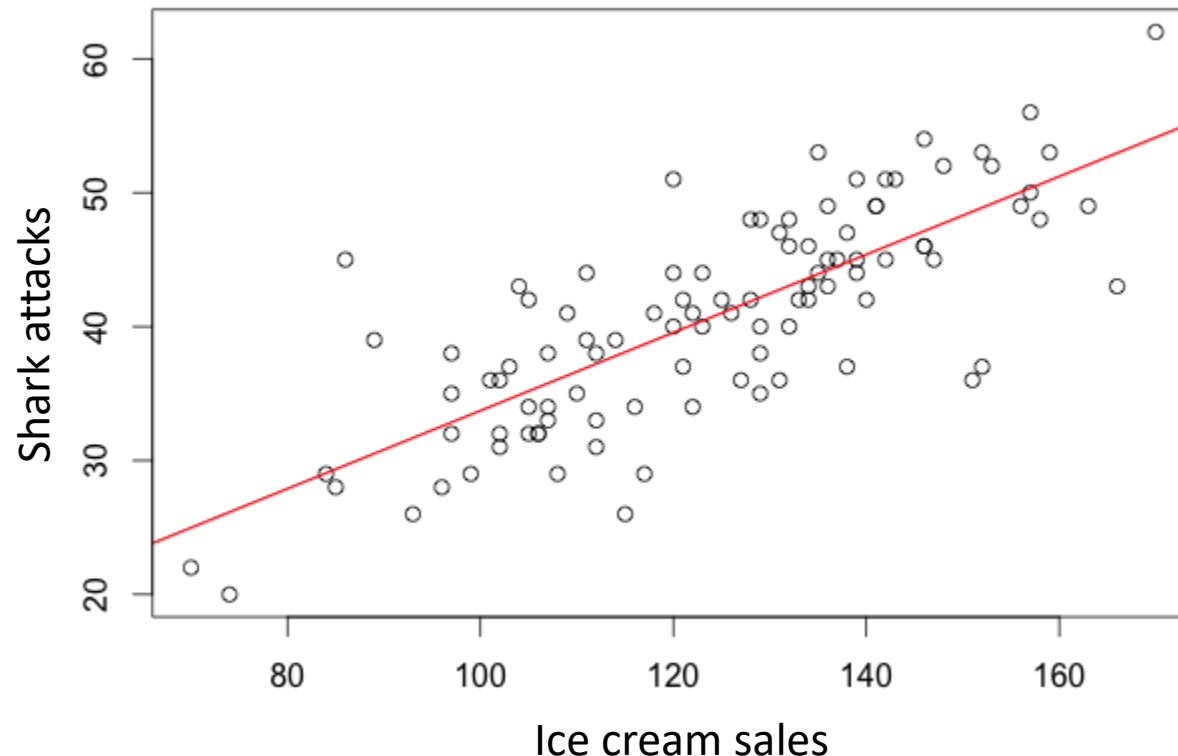
(c) Alkalinity vs acidity



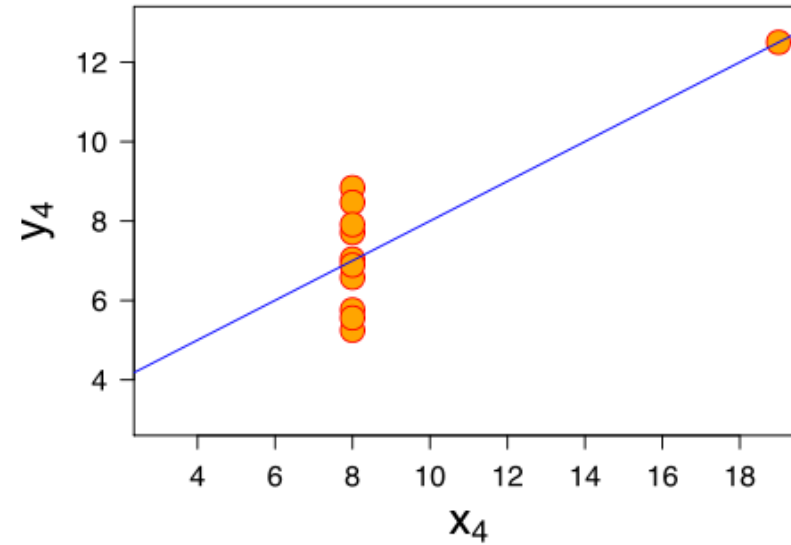
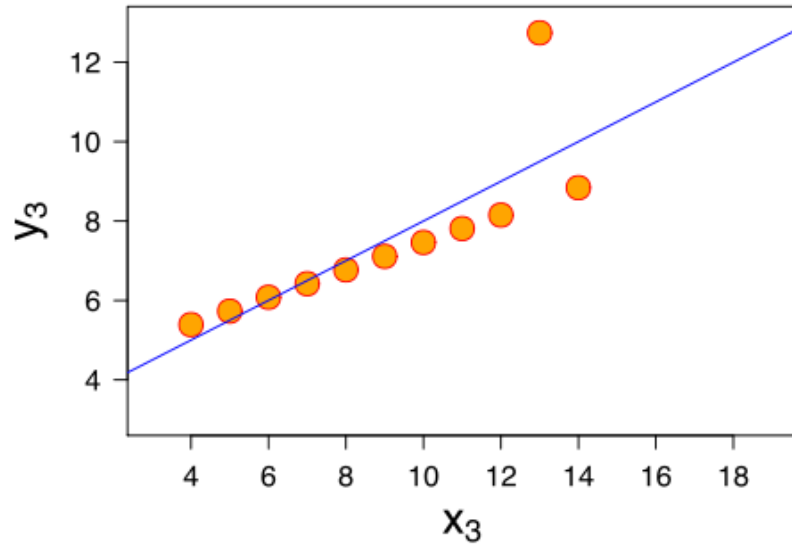
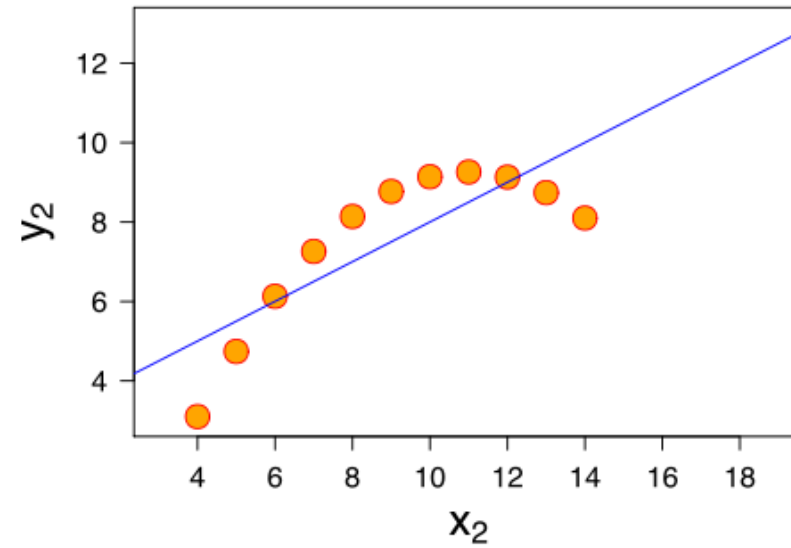
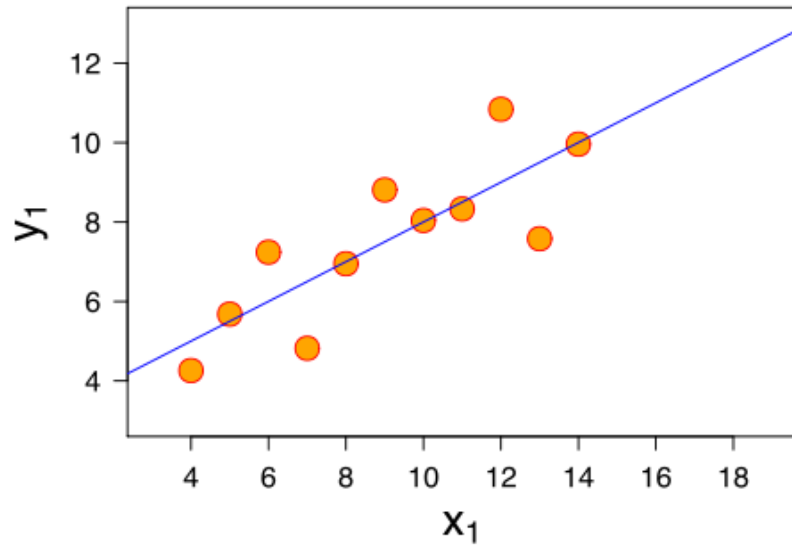
(d) Average vs standardized mercury levels

# Correlation caution #1

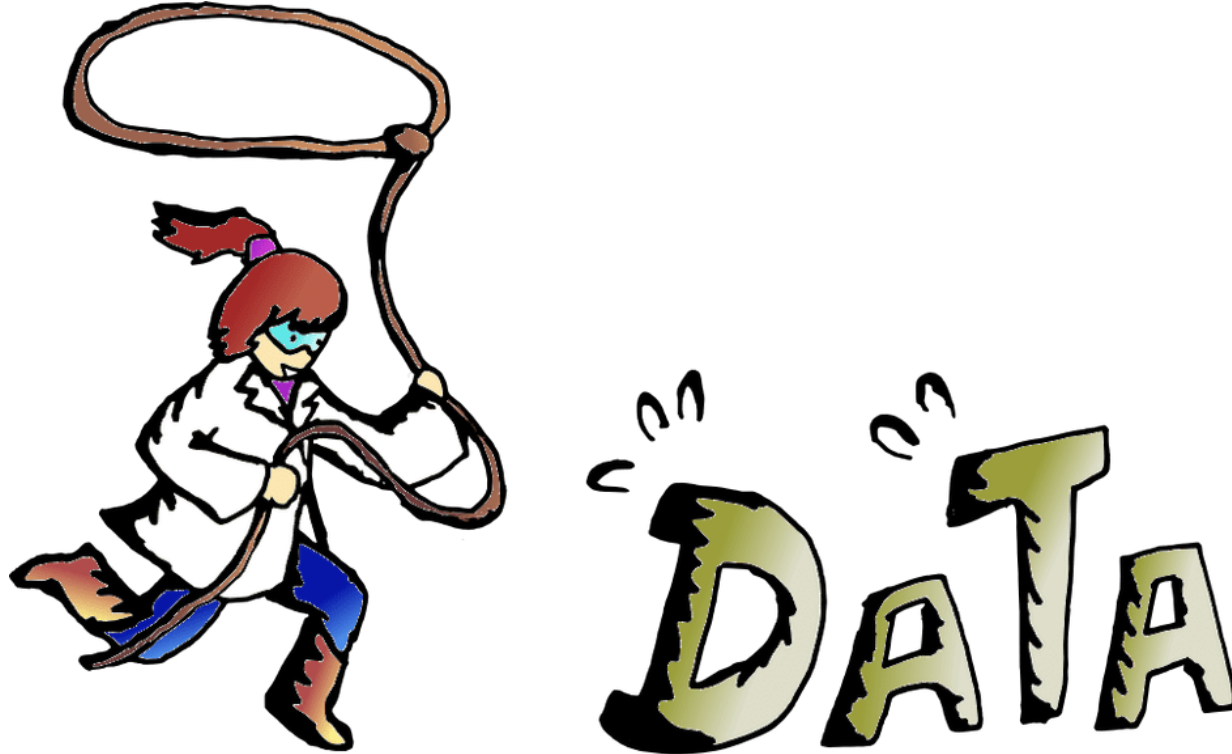
A strong positive or negative correlation does not (necessarily) imply a cause and effect relationship between two variables



# Anscombe's quartet ( $r = 0.81$ )



# Data wrangling/transformation using dplyr



The tidyverse and dplyr

# The 'tidyverse'

The tidyverse is set of R packages that operate 'tidy data'

- i.e., that operate on data frames (or tibbles)

Tidy data is data where:

- Each variable must have its own column
- Each observation must have its own row
- Each value must have its own cell



country	year	cases	population
Afghanistan	1999	745	15987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272412272
China	2000	216766	1280425583

variables

country	year	cases	population
Afghanistan	1999	745	15987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272412272
China	2000	216766	1280425583

observations

country	year	cases	population
Afghanistan	1999	745	15987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272412272
China	2000	216766	1280425583

values

# Messy data...

## What would be an example of data that is not tidy?

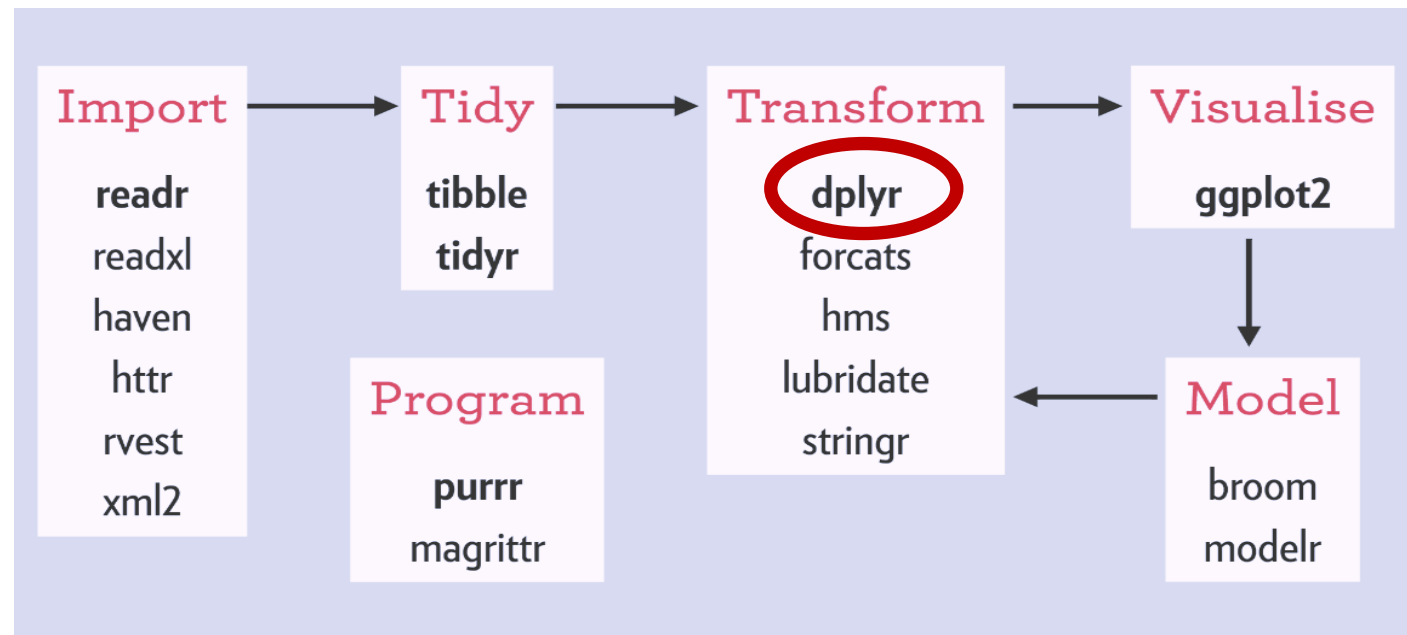
[illegible]



# The 'tidyverse'

The tidyverse is a set of packages share a common design philosophy

- Most written by Hadley Wickham



# dplyr: A grammar for data wrangling

**Grammar:** a set of components that can be combined to achieve a goal

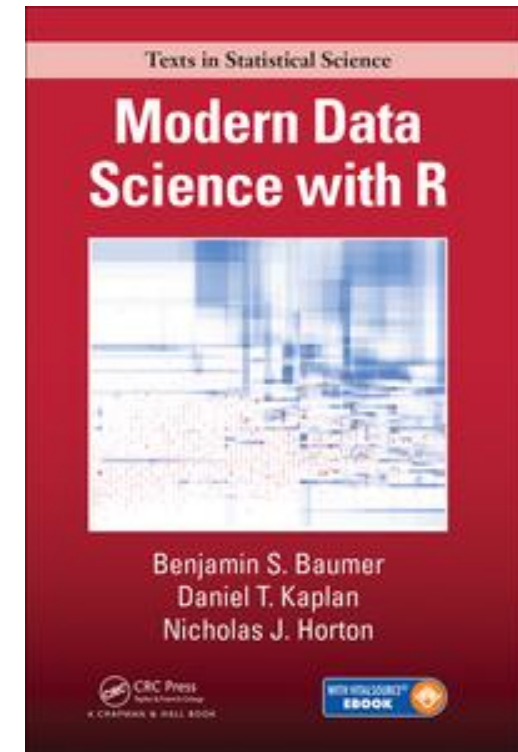
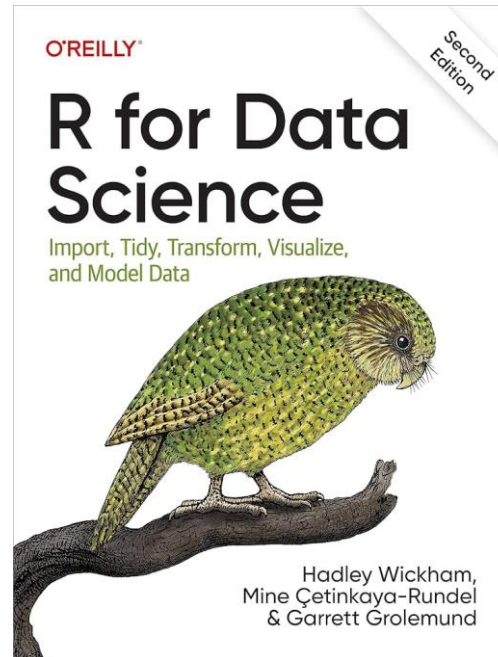
**dplyr** is a package that has a set of verbs that are useful for transformations data:

1. `filter()`
2. `select()`
3. `mutate()`
4. `arrange()`
5. `group_by()`
6. `summarize()`

All these function **take a data frame** and other arguments and **return a data frame**

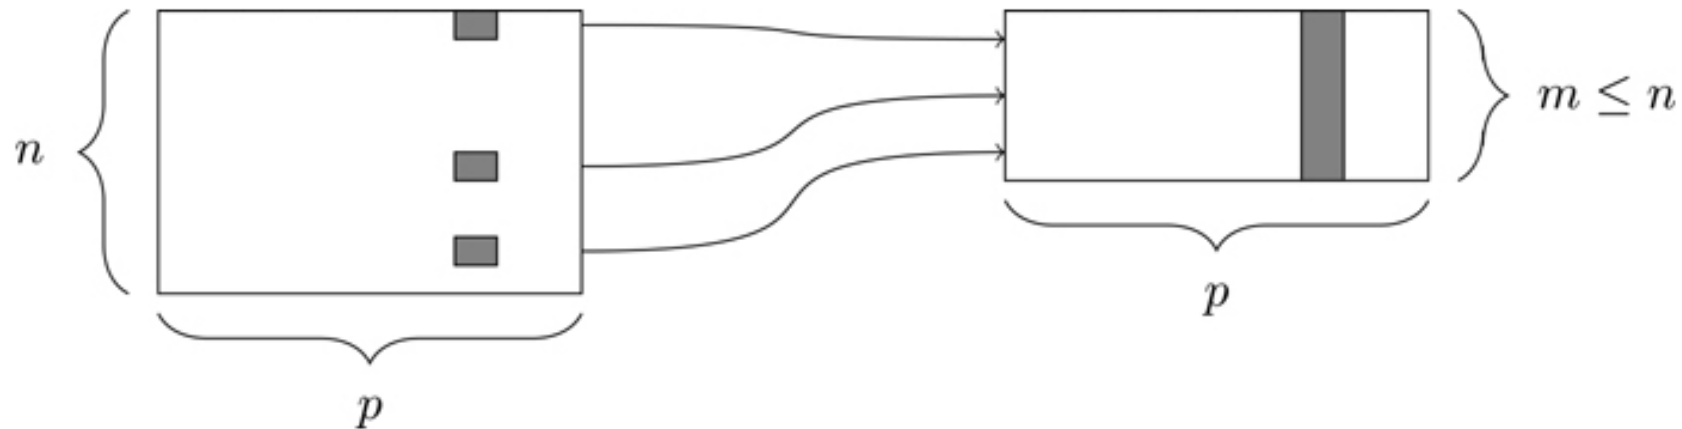
```
> library(dplyr) # load the dplyr package
```

# Quick overview of the dplyr functions



# 1. filter()

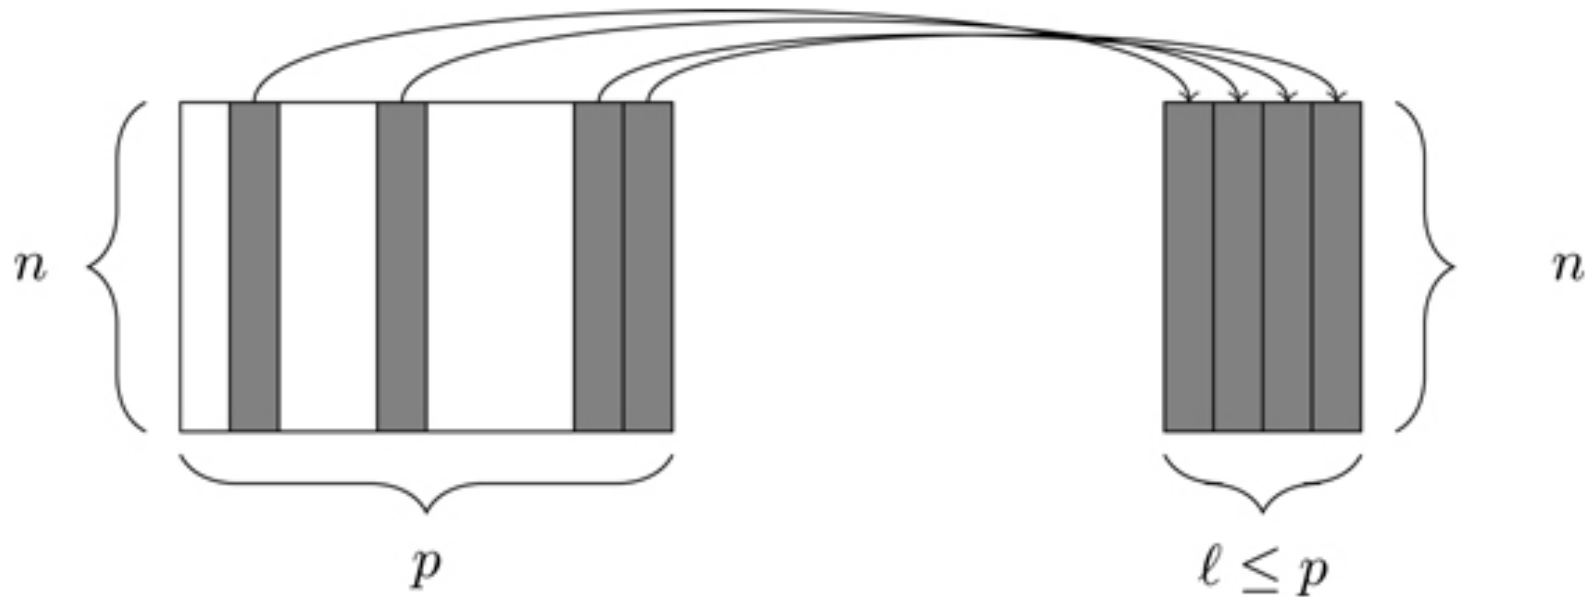
The `filter()` function allows you to select a subset of rows in data frame



```
filter(profiles, height == 77)
```

## 2. select()

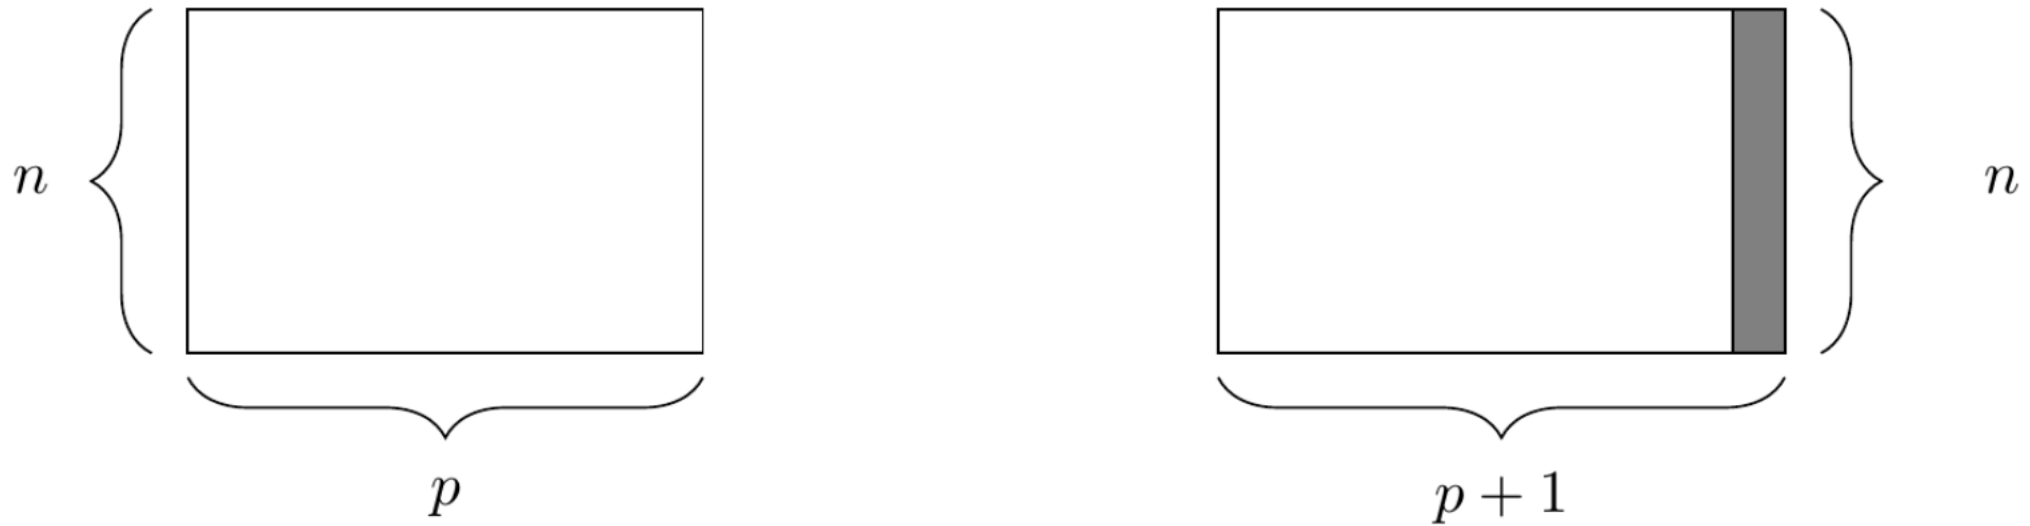
The `select()` function allows you to select a subset of columns



```
select(profiles, age, height)
```

### 3. mutate()

The `mutate()` function allows you to create new columns that are functions of existing columns

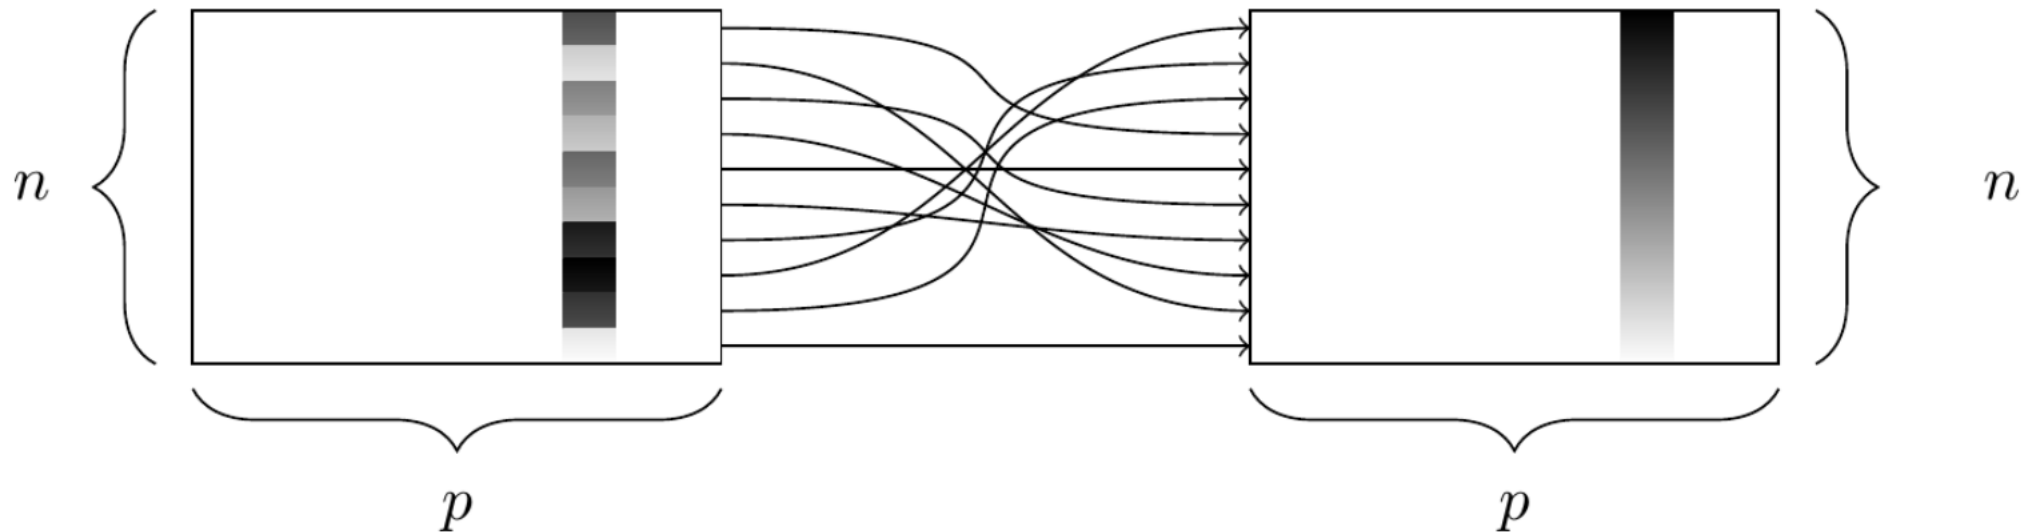


```
mutate(profiles, height_feet = height/12)
```

## 4. arrange()

The `arrange()` function arranges the rows based values in a column

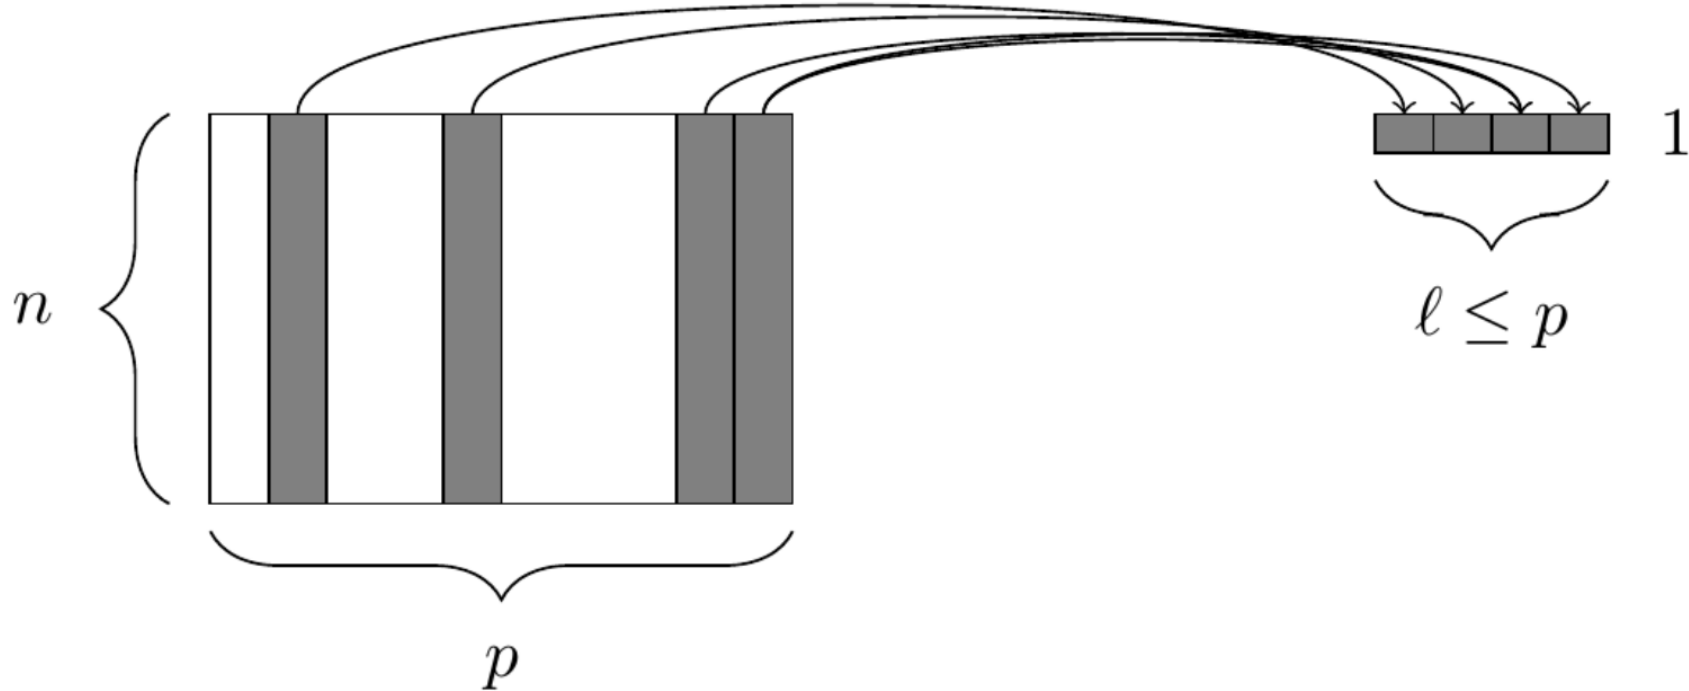
- `arrange(desc())` arranges from largest to smallest



`arrange(profiles, desc(height))`

## 5. summarize()

The `summarize()` function reduces values in many rows into single values



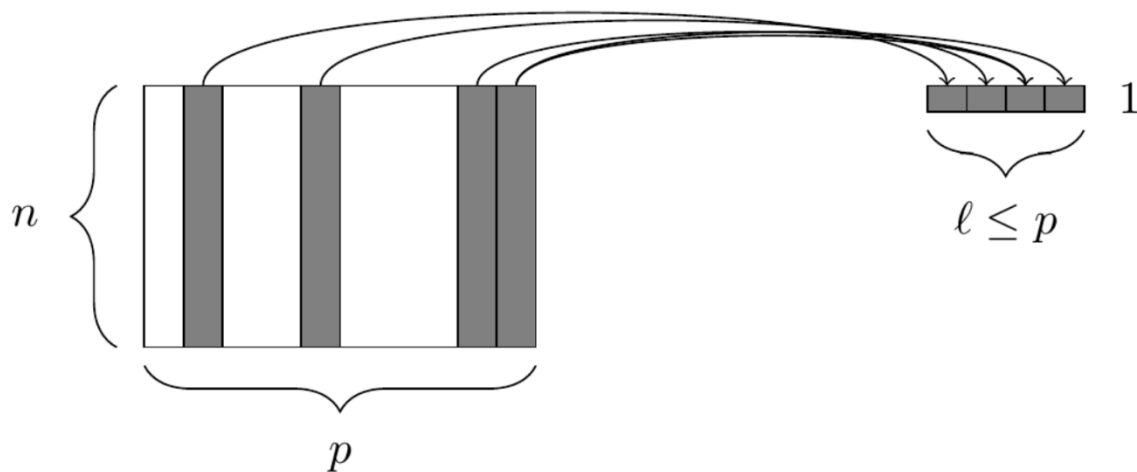
```
summarize(mean_age = mean(age))
```



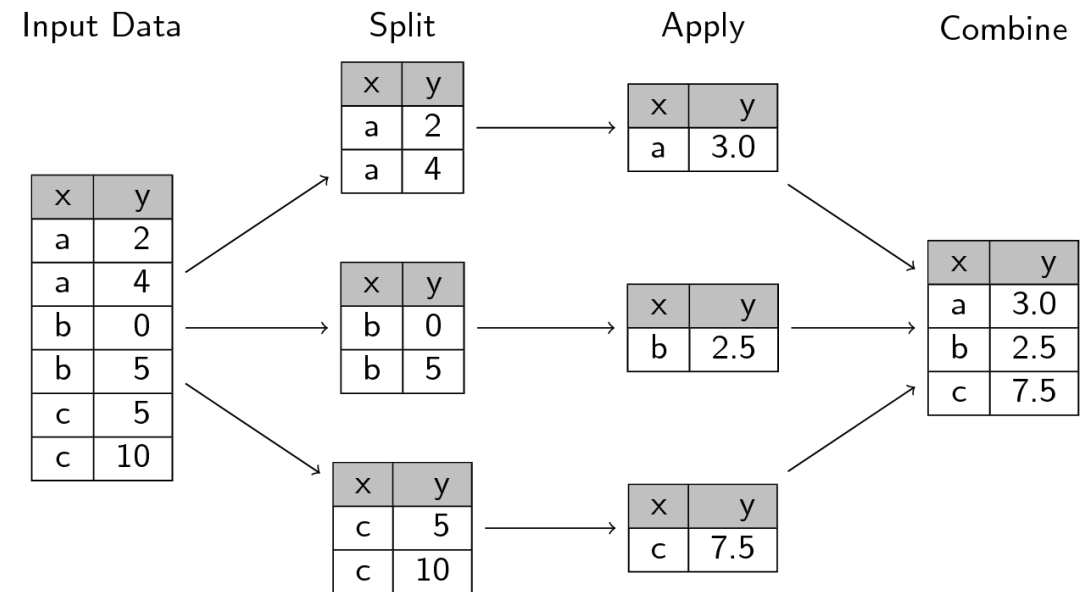
## 6. The group\_by() function

The `group_by()` function groups variables for future operations

- It works in conjunction with `summarize()` and `mutate()` to do **split, apply, combine**

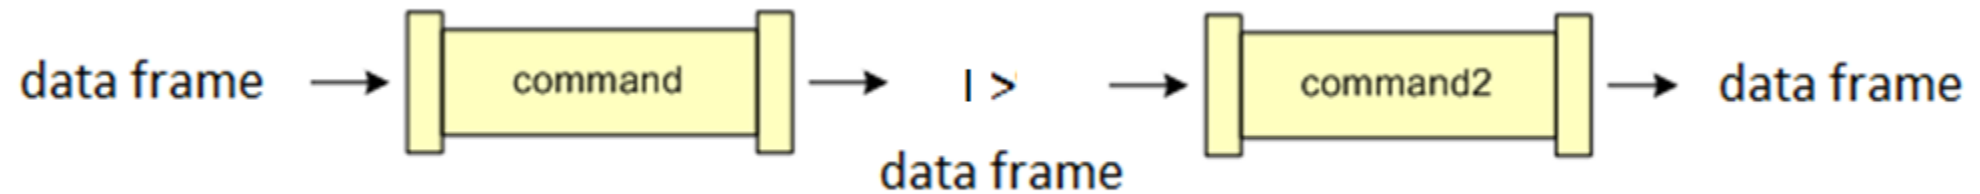


`group_by(profiles, sex)`



# The pipe operator

The pipe operator `|>` allows us to chain commands together



`profiles|>`

`group_by(sex) |>`

`summarize(mean_age = mean(age))`



Let's try it out!